

Linear Regression

Dr. Kuppusamy .P
Associate Professor / SCOPE

Linear Regression

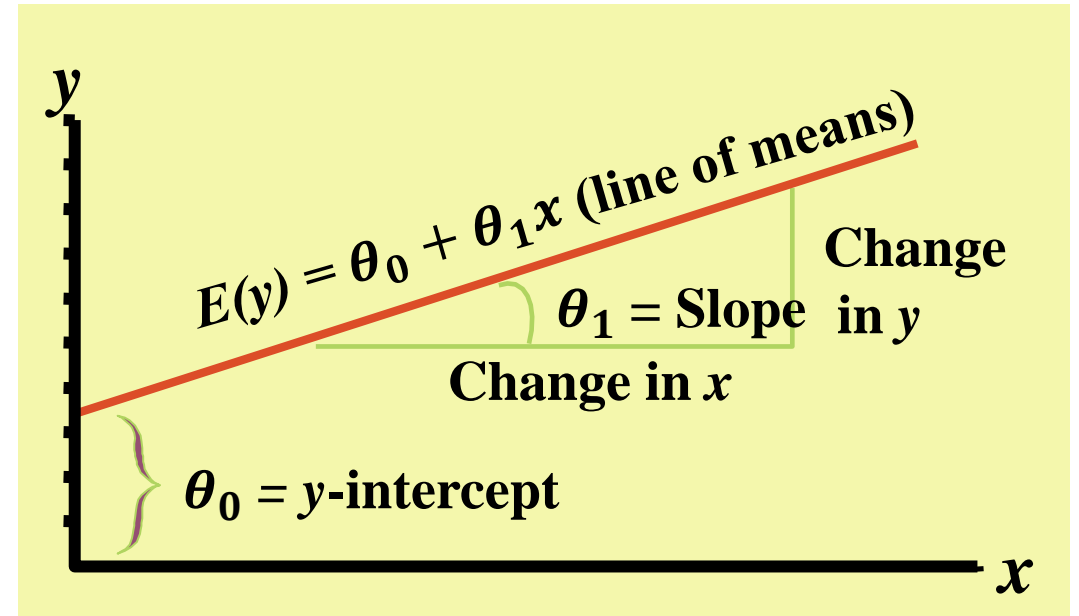
- Linear Regression analysis is a statistical (quantitative analysis) tool
- Predictive modeling method to investigate the mathematical relationship between an independent variable (predictor – x) and **continuous** dependent variable (outcome – y) and
- Predictor shows the changes in Dependent variable (y axis) to the changes in explanatory variables in X axis.
- It uses current information about a phenomenon to predict its future behavior.
- Involves the graphical lines over a set of data points that most closely towards all shape of the data.
- When the data form a set of pairs of numbers, it is interpreted as the observed values of an independent (or predictor) variable X and a dependent (or response) variable Y .

x	y
1	3
2	4
3	2
4	4
5	5

Data model in Simple Linear Regression

- Data is modelled using a **straight line** with continuous variable
- Relationship between variables is a linear function.
- Linear equation representation with single feature is given:



$$\underset{\substack{\text{Dependent (Response)} \\ \text{Variable}}}{y} = \underset{\substack{\text{Population} \\ \text{y-intercept}}}{\theta_0} + \underset{\substack{\text{Population Slope}}}{\theta_1} \underset{\substack{\text{Independent} \\ \text{(Explanatory) Variable}}}{x} + \underset{\text{Random Error}}{\varepsilon}$$





- It is reduced into
$$y' = h_{\theta}(x) = \theta^T \cdot X$$

Types of Relationships



Strong relationships



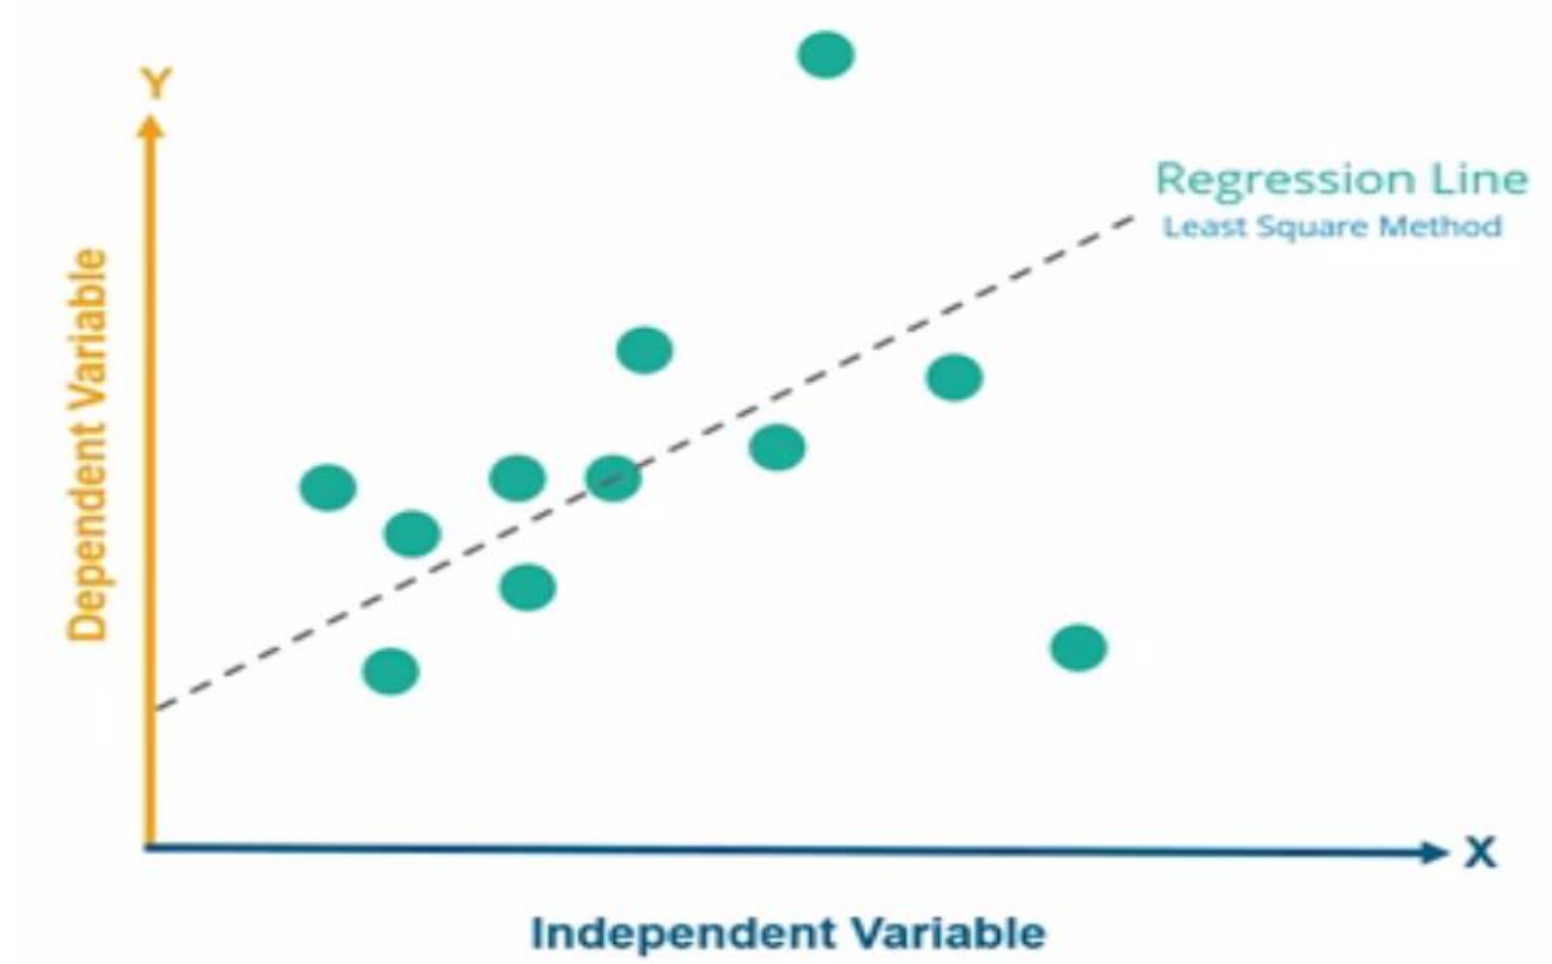
Weak relationships



No relationship




Plot the graph using x and actual y (target) values



Random Error Identification

- Random Error $\varepsilon = \text{Estimated Value } (y_i') - \text{Actual Value } (y_i)$



Minimize the Random Error using Cost Function **Least squared (LSE) Method**


- Reduce the distance between estimated (predicted) and actual (target) value
- Find the best fit of the line using **least square method**.

Least Squares Method to Minimize the Error

- **Best fit** means difference between actual y values and predicted y values are a **minimum**
 - *But* positive differences off-set negative

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n \hat{\varepsilon}_i^2$$

- $y' = h_{\theta}(x) = \theta^T \cdot X$
- Least Squares minimizes the **Sum of the Squared** Differences (SSE)



Minimize the Random Error using Cost Function

- Normal Equation to calculate the weights

$$\theta = (X^T \cdot X)^{-1} \cdot (X^T \cdot y)$$


- Compute the inverse of the matrix $(X^T \cdot X)^{-1}$ and $X^T \cdot y$ (dot) matrix vector multiplication.

Computational complexity

- In least squares regression, Assume N training examples and C features. It consumes
 - $O(C^2N)$ to multiply X^T by X
 - $O(CN)$ to multiply X^T by Y
 - $O(N^3)$ i.e., $C=N$ to compute the LU (or Cholesky) factorization of $X^T X$ and use that to compute the product $(X^T \cdot X)^{-1} \cdot (X^T \cdot y)$
- The computational complexity with the normal equation about $O(n^{2.4})$ to $O(n^3)$.
- For large set (100 000) of features, normal equation computation is slow.
- However linear equation handles large training sets well. So, model can fit into memory.
- Once the model is trained the new predictions will be fast.

Least Squares Graphically

$$\text{LS minimizes } \sum_{i=1}^n \hat{\varepsilon}_i^2 = \hat{\varepsilon}_1^2 + \hat{\varepsilon}_2^2 + \hat{\varepsilon}_3^2 + \hat{\varepsilon}_4^2$$



Minimize the Random Error using **Mean squared Error (MSE)**


- Reduce the distance between estimated (predicted) and actual (target) value
- Find the best fit of the line using **Mean squared Error (MSE) method**.

Mean Squared Error to Minimize the Error

- **Best fit** means difference between actual y values and predicted y values are a **minimum**
 - *But* positive differences off-set negative

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m ((h_{\theta}(x_i) - y_i))^2$$

- $y' = h_{\theta}(x) = \theta^T \cdot X$



Multi-variate Linear Regression

- **Multi-variate Linear Regression**
- Linear equation representation with **multiple features**


$$y = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

- It is reduced into

$$y' = h_{\theta}(x) = \theta^T \cdot X$$

$$\text{Cost Function } J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)})^2$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left(\left(\sum_{j=0}^n \theta_j x_j^{(i)} \right) - y^{(i)} \right)^2$$



Gradient Descent

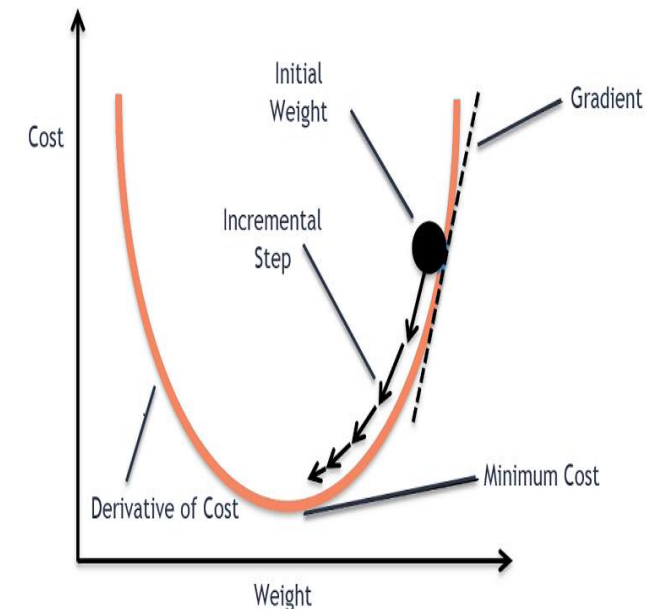

- Gradient descent is an optimization algorithm tweak the parameters iteratively in order to find the optimal cost function.

Idea of Gradient Descent

- Let consider you are on the mountain peak in the dense fog, and you want to go down. You can only feel the slope of the ground below your feet. A good strategy to reach the bottom of the valley quickly is to go downhill in the direction of the steepest slope.
- Gradient descent does the above context. It measures the local gradient of the error function with regards to the parameter vector θ , and it goes in the direction of descending gradient. Once the gradient is close to zero or zero, you have reached a **minimum**.
- Initialize θ as random values and then improve one step at a time to decrease cost function.
- The size of the steps determined by the learning rate parameter. If choose learning rate is too small, the algorithm takes many iterations to achieve the minimum.
- Convergence theorem for gradient descent is $\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta} J(\theta_0, \theta_1); j = 0 \text{ and } j = 1$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m ((h_{\theta}(x_i) - y_i))^2$$

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$



Learning Parameter in Gradient Descent

- The size of the steps determined by the learning rate parameter.

For Small Learning rate:



- If choose learning rate is too small, the algorithm takes many iterations to achieve the minimum.

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m ((h_{\theta}(x_i) - y_i)x_i)$$


For High Learning rate:

- If choose learning rate is too high, the algorithm may jump across the valley in finding solution even higher than before.



Learning Parameter in Gradient Descent


- However, all functions does not contain regular shape.
- Few functions may contain local minimums which are not as good as the global minimum.



MSE cost function in Linear Regression

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m ((h_{\theta}(x_i) - y_i))^2$$

- But MSE cost function for a Linear Regression model contains convex function i.e., if two points selected on the curve, the line segment joining them never crosses the curve.
- This implies that there are no local minima, just one global minimum. It is also a continuous function with a slope that never changes abruptly.
- It leads that GD guarantees to approach close to the global minimum.




Case Study – Simple Linear Regression (i.e., **Single feature x_1**)



- Let consider x and y values and mark in scatter plot
- Calculate slope **m** for the **regression line** $y=mx+c$.

- $$m = \frac{\sum(x-\bar{x})(y-\bar{y})}{\sum(x-\bar{x})^2}$$


- \bar{x} is mean of all x
- \bar{y} is mean of all y.



- Find mean of x, and mean of y





- Find the coefficients m and c in the straight line $y = mx + c$




$$m = \frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2}$$

- Find $x - \bar{x}$, and $y - \bar{y}$




- Find slope m and y-intercept c .



x	y	$x - \bar{x}$	$y - \bar{y}$	$(x - \bar{x})^2$	$(x - \bar{x})(y - \bar{y})$
1	3	-2	-0.6	4	1.2
2	4	-1	0.4	1	-0.4
3	2	0	-1.6	0	0
4	4	1	0.4	1	0.4
5	5	2	1.4	4	2.8
3 3.6		$\Sigma = 10$		$\Sigma = 4$	


$$m = \frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2} = \frac{4}{10}$$

Plot the x, y values in the graph $x = \{1, 2, 3, 4, 5\}$ $y = \{3, 4, 2, 4, 5\}$



Plot the regression line using **estimated $y = (mx+c)$ values**

$$x = \{1, 2, 3, 4, 5\} \quad y = \{2.8, 3.2, 3.6, 4, 4.4\}$$




$$\begin{aligned} m &= 0.4 \\ c &= 2.4 \\ y &= 0.4x + 2.4 \end{aligned}$$

For given $m = 0.4$ & $c = 2.4$, lets
predict values for y for $x = \{1, 2, 3, 4, 5\}$

Estimated y values

$$\begin{aligned} y &= 0.4 \times 1 + 2.4 = 2.8 \\ y &= 0.4 \times 2 + 2.4 = 3.2 \\ y &= 0.4 \times 3 + 2.4 = 3.6 \\ y &= 0.4 \times 4 + 2.4 = 4.0 \\ y &= 0.4 \times 5 + 2.4 = 4.4 \end{aligned}$$

Find the Error ε



$$\sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n \hat{\varepsilon}_i^2$$

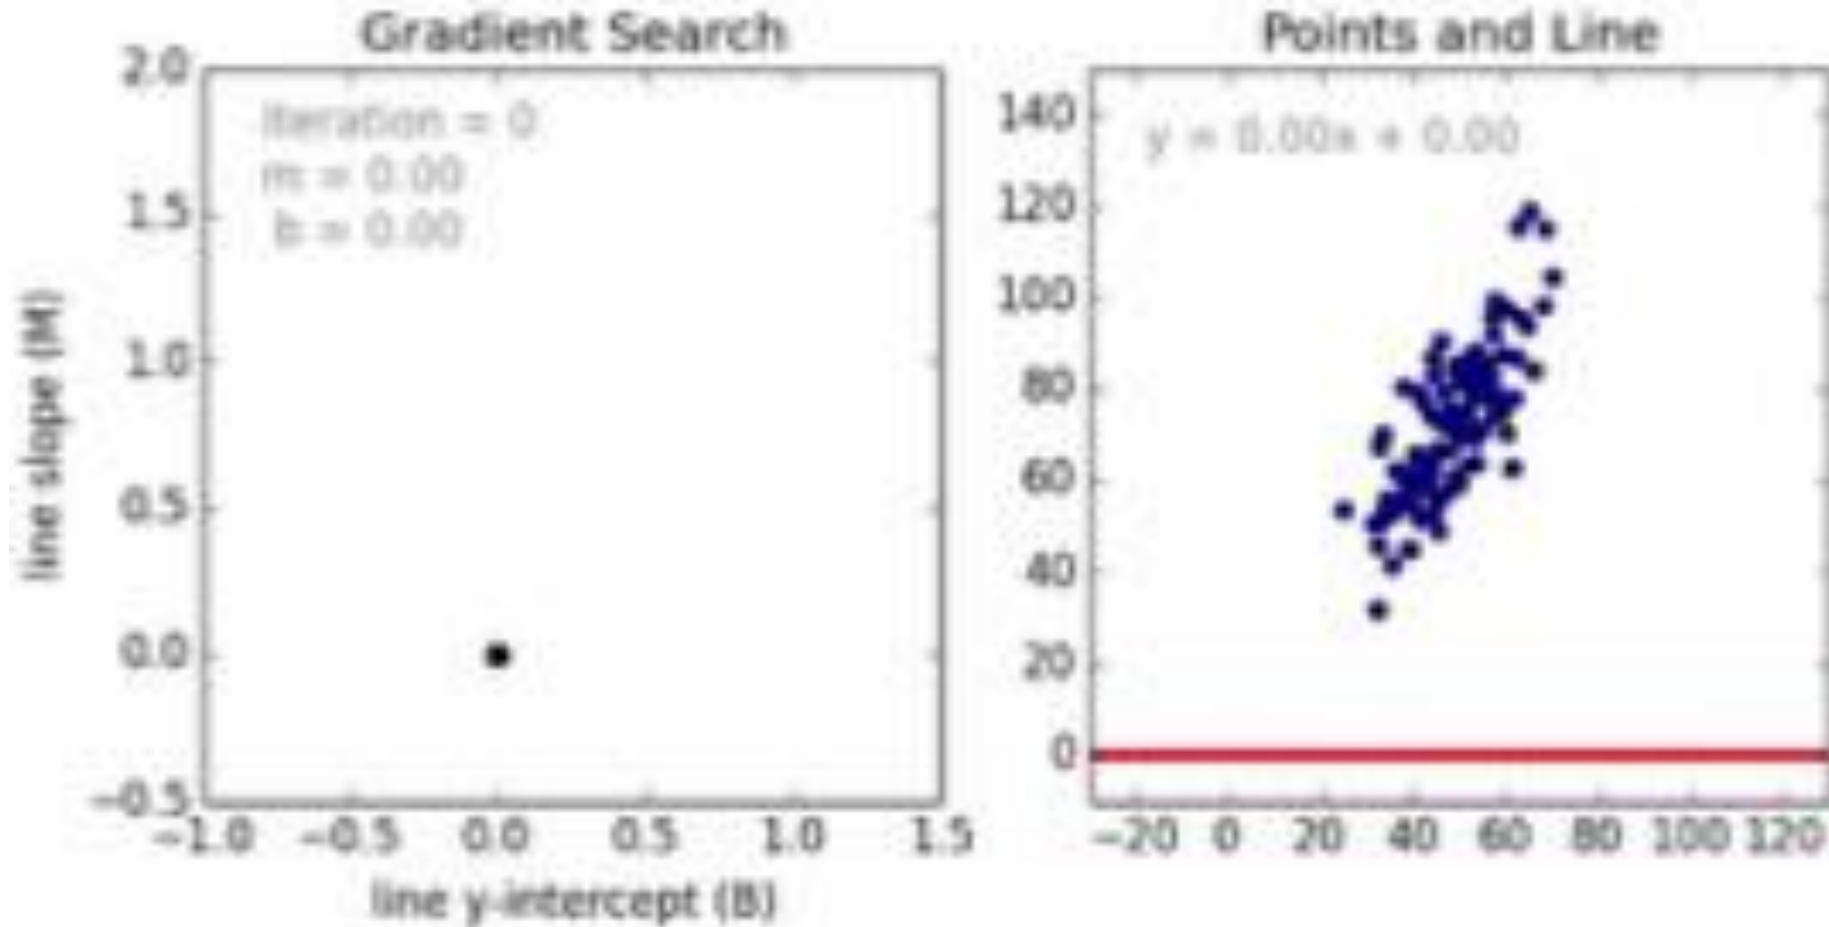
Distance between actual
& predicted value

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

- Mean Square Error minimizes the error in the linear regression.
- Regression Line with least error is the 'best fit' line

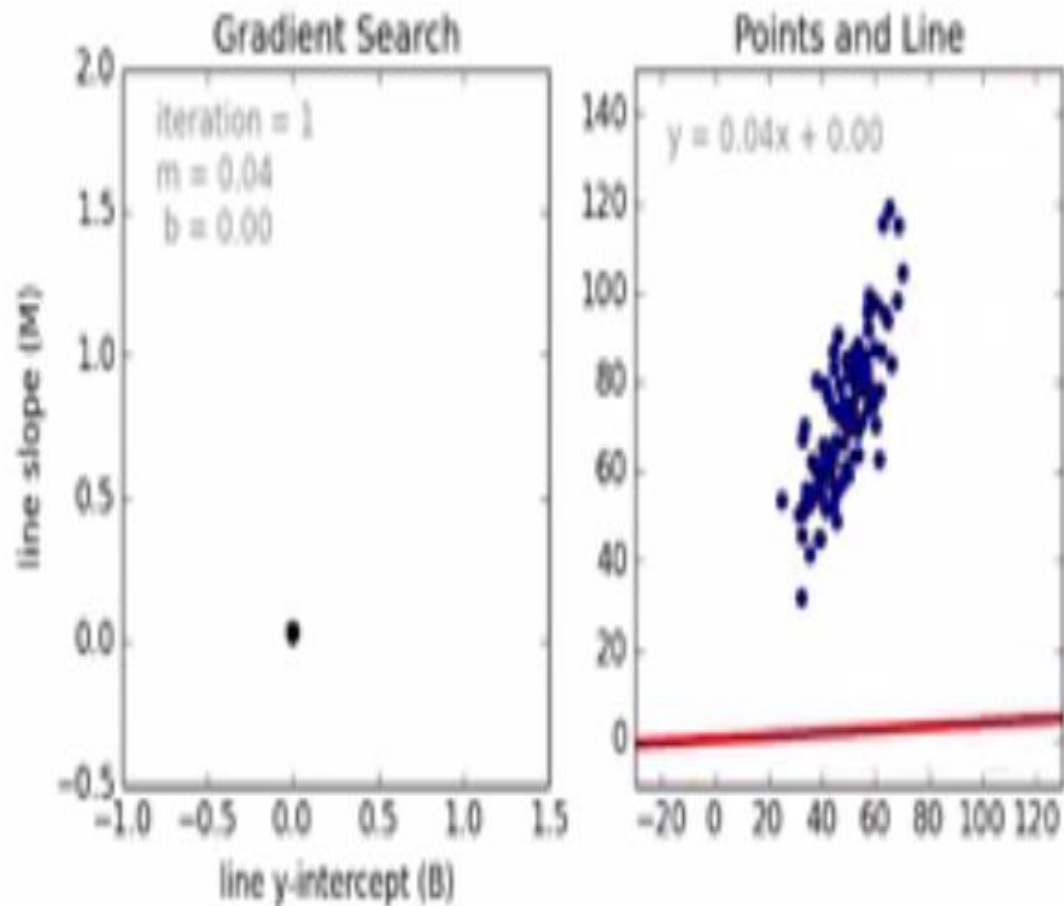
How would you draw a line through the points in real time?

- Initial values (iteration 0) for slope $m = 0$ and y-intercept $b = 0$




How would you draw a line through the points?

iteration 1, slope $m = 0.04$ and y-intercept $b = 0$




iteration 20, slope $m = 0.59$ and y-intercept $b = 0.01$




Determine which line 'fits best' in 100 iterations

iteration 47, slope $m = 1.03$ and y-intercept $b = 0.02$



iteration 99, slope $m = 1.36$ and y-intercept $b = 0.03$



Applications of Regression

- Forecasting an effect
- Trend forecasting and sales estimates
- Analyze the impact of price changes
- Insurance domain

References

1. Tom M. Mitchell, Machine Learning, McGraw Hill , 2017.
2. EthemAlpaydin, Introduction to Machine Learning (Adaptive Computation and Machine Learning), The MIT Press, 2017.
3. Wikipedia